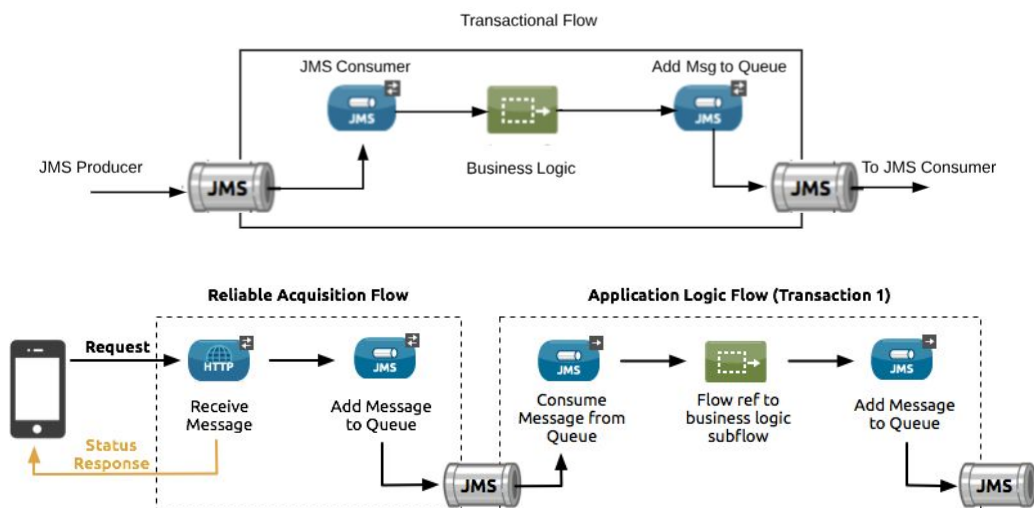
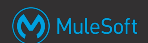




Module 12: Designing for Reliability Goals

Goal



At the end of this module, you should be able to



- Distinguish between competing non-functional requirements
- Clarify and validate reliability goals for a scenario
- Design Mule applications and their deployments to meet reliability goals
- Identify reliability pattern for mule application and their deployments

Balancing tradeoffs to meet non-functional requirements



Identifying non-functional requirements



- Applications can be designed and tuned for various opposing goals
 - High availability
 - Reliability
 - Performance
 - Fast response time/low latency
 - High throughput
 - Capacity (number of concurrently processed messages)
 - Security
- Performance requirements, service level agreements (SLAs), and other business goals should be clearly agreed upon and documented
 - They are often opposing goals

Considering tradeoffs to meet opposing non-functional requirements



- Cost
- Reliability
 - Business implications of duplicate message processing vs. lost messages
- Competing performance SLAs
 - Fast response time/low latency
 - High throughput
 - Capacity (number of concurrently processed messages)
- Transactional exactly once requirements
 - Latency and cost vs. business risk

Defining and achieving reliability goals



Defining and achieving reliability goals



- **Reliability** aspires to have zero message/data loss after a Mule application stops or crashes
- Various reliability patterns can be implemented to achieve reliability goals for synchronous and asynchronous flows
- Reliability in Mule applications can be achieved using
 - Until Successful scope
 - Reconnection strategies
 - Redelivery policy
 - NS:RETRY_EXHAUSTED exception scope
 - Transactions (covered in module 11 - Transaction)

Achieving reliability goals with Mule components



Achieving reliability using an Until Successful scope




- The **Until Successful** scope repeatedly triggers the scope's components (including flow references) **until they all succeed or until a maximum number of retries is exceeded**
- The scope provides option to control the max number of retries and the interval between retries
- The scope can execute any sequence of processors that may fail for whatever reason and may succeed upon retry


```
<until-successful maxRetries="5"
  millisBetweenRetries="1000">
  <!-- One or more processors here -->
</until-successful>
```



Achieving reliability reconnection strategies



- **System errors** are thrown when a connection to an external system (DB, message broker, etc) fails
- To retry after connection failures, Mule connectors can set a **reconnection strategy**
 - Set for a connector (in the Global Elements Properties) or for a specific connector operation (in the Properties view)




All contents © MuleSoft Inc.

11

Achieving reliability using a redelivery policy



- A **redelivery policy** is configured on inbound connectors, such as the JMS connector, to specify the number of redeliveries before discarding the message
 - 0 means no redelivery
 - -1 means infinite redeliveries



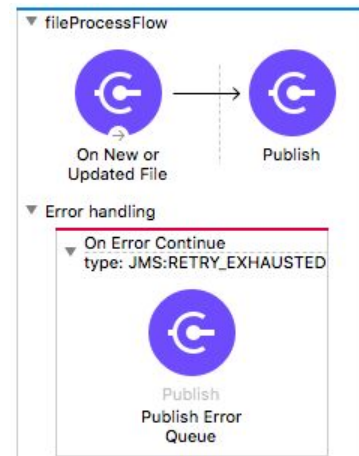
All contents © MuleSoft Inc.

13

Achieving reliability using MuleSoft components - RETRY_EXHAUSTED exception scope



- Before discarding the message after the number of redeliveries attempted, the connectors raises an exception of type `RETRY_EXHAUSTED`
- An error scope can handle the `RETRY_EXHAUSTED` error with logic required to handle the error, so the event is not lost



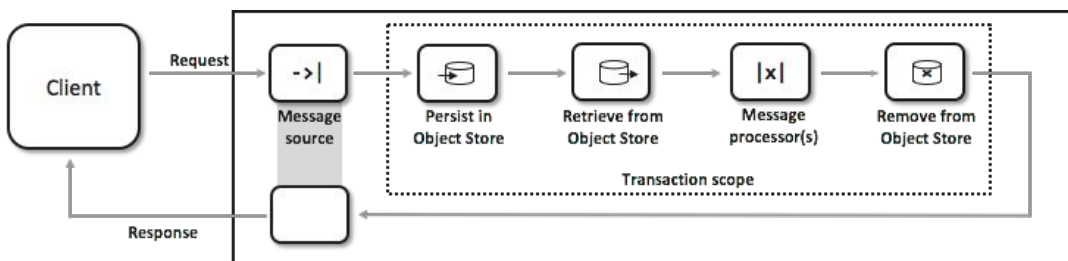
All contents © MuleSoft Inc.

14

Achieving reliability for transactional systems



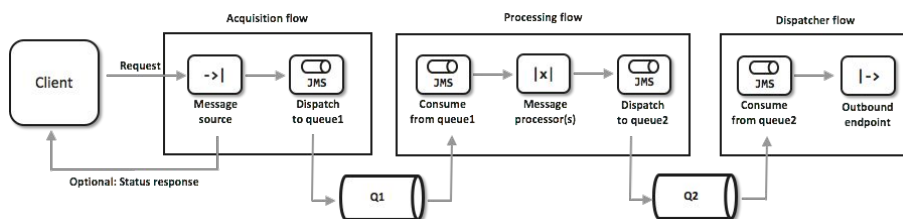
- Zero message/data loss for transactional systems is achieved using a **transaction**
- Message Ack is another way to achieve zero message loss
- Message persistence for application downtime/crash is required
- More details are provided in the persistence module



All contents © MuleSoft Inc.

16

- Zero message/data loss for non-transactional systems is achieved using a **reliability pattern**
- Splits processing between an **acquisition** flow and a **processing** flow
- The flows do not call each other directly, but use **persisted queues**
- In case of application failure, messages/data are still available in those queues

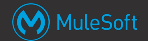


17

Understanding the two flows in this reliability pattern for non-transactional systems

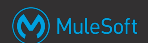
- The reliability pattern consists of flows with specific responsibilities
 - The **Acquisition flow**
 - Receives incoming messages then dispatches them to a persisted processing queue
 - In case of failure, processes the message until redelivery exhausted, and if that fails, then dispatches the message to a persistent error queue
 - The **Processing flow**
 - Process messages from the processing queue then dispatches them to another persisted queue
 - In case of failure, processes the message until redelivery exhausted, and if that fails, then dispatches the message to a persisted error queue

Example: The acquisition flow of a reliability pattern for a non-transactional systems

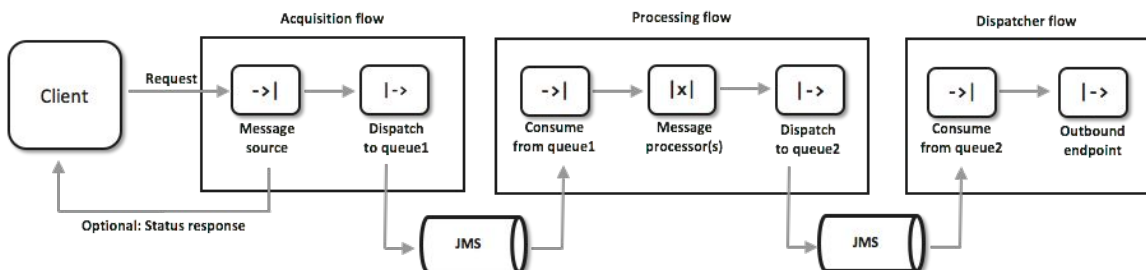


```
<!-- File listener configure with redelivery count of n for in case of failure -->
<file:listener>
    <redelivery-policy maxRedeliveryCount="n" />
</file:listener>
<!-- file is persisted in VM queue -->
<vm:publish doc:name="Publish" queueName="file"/>
<!-- In case of any error occurs, the read will be rolled back and the message processed
until maxRedeliveryCount and finally persisted in error queue -->
<on-error-continue type="REDELIVERY_EXHAUSTED">
    <vm:publish queueName="fileerror"/>
</on-error-continue>
```

Achieving reliability for non-transactional systems



- The processing flow must read the message queue transactionally
- Queues can be persistent VM queues or JMS queues
- A redelivery policy is set on event sources in both flows
- REDELIVERY_EXHAUSTED type errors are handled in both the acquisition and processing flows



Other ways to achieve reliability with Mule applications



- In an earlier module you already learned about how to store state in a Mule application
 - Object Store
 - Persistent queues
 - File persistence
 - Database connector
 - Caches
 - Other external systems
- Each option has tradeoffs between reliability goals vs. performance vs. cost

Summary



- Application level persistence (file, OSv2) achieves different levels of reliability depending on the runtime planes and configurations
- Reliability patterns for synchronous and asynchronous flows ensure reliability in the Mule application through persistent queues across various connectors