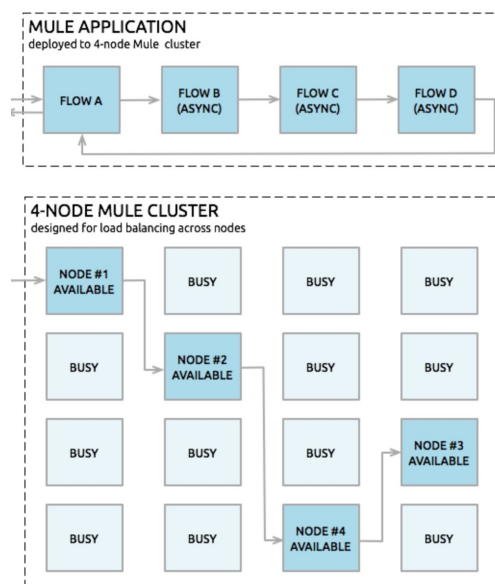


Module 13: Designing for High Availability Goals

Goal



At the end of this module, you should be able to



- Clarify HA goals for Mule applications
- Balance HA goals with reliability and performance goals
- Identify ways to achieve high availability (HA) using Anypoint Platform, in CloudHub and on-premises
- Describe how clustering and load balancing works
- Identify HA aware connectors and their design tradeoffs

Achieving high availability (HA) goals using multiple Mule runtimes



Distinguishing between high availability (HA) vs. disaster recovery (DR) goals





- High availability (HA)
 - How to keep the overall system operational when a system component fails
 - Usually achieved with multiple levels of fault tolerance and/or load balancing
 - In CloudHub, you can deploy a Mule application to multiple CloudHub workers
- Disaster recovery (DR)
 - How to restore a system to a previous acceptable state after a natural or man-made disaster
 - Must compensate for failure of entire sub-systems/regions/zones/data centers, and hence there are a large number of system components that failed
- Read details in the Mule User Guide

<https://docs.mulesoft.com/mule-user-guide/v/3.9/hadr-guide>

Making Mule applications highly available using Mule runtimes



- High availability can be achieved by **horizontally scaling** to multiple Mule runtimes
 - Process on multiple concurrent physical machines/VMs, often distributed across networks
- HA goals can be met via **load balancing** and/or **clustering**
 - Clustering of Mule runtimes uses an active-active model of node
 - Load distributes across the active nodes




Comparing Mule runtime server groups vs. clusters




- Runtime Manager can define server groups or clusters
- A server group is just an administrative grouping of isolated Mule runtimes
- A cluster provides additional guarantees to prevent contention between the Mule runtimes
 - File access by File based transports
 - All JMS topic subscribers connect to the same topic, resulting in duplicate processing
 - JMS request/response queues might send the response to a different Mule runtime, causing uncorrelated response processing or other failures
 - Salesforce streaming API will fail because the API only supports a single consumer

<https://docs.mulesoft.com/mule-runtime/4.1/mule-high-availability-ha-clusters>

Unclustered load balancing for HA and performance in customer-hosted runtime planes



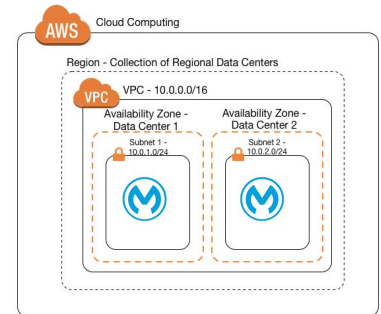
- Multiple Mule runtimes are configured to run the same Mule application(s)
- Does not share or synchronize data between Mule runtimes
 - Could potentially lead to processing duplicate or lost messages
 - But using a shared database, message broker, or other system is possible
- Messages must be distributed or load balanced
 - Requires third party products



Achieving HA in CloudHub using load balancing between multiple CloudHub workers



- HA can be achieved using multiple CloudHub workers (>1)
 - Easy to configure
 - Workers do not share any memory
 - Workers can use an external system for state management
- Each worker is created in a **different availability zone** in the **same AWS region**
- Mule application data can be stored and shared between CloudHub workers in the Anypoint Object Store (OSv2)
- VM queues can be changed to persistent or non-persistent in Runtime Manager
 - Without changing any configuration or properties of the deployed Mule app



All contents © MuleSoft Inc.

9

Achieving HA using customer-hosted clusters



- A cluster is a set of customer-hosted Mule runtimes that act as a unit
 - Servers in a cluster communicate and share information through a distributed shared memory grid
 - Selected data is replicated across memory in different physical machines
 - Cluster nodes can be configured to be more reliable so they also copy data to disk or external storage
- Nodes are aware of each other
- All cluster nodes work in active-active mode and there is a primary node where Schedulers, JMS listeners, etc. run
- To manage peak loads, nodes can be added to / subtracted from a cluster
- Only available for customer-hosted Mule runtimes

All contents © MuleSoft Inc.

10

How clusters survive failure of a Mule runtime



- If one Mule runtime (node) fails, outstanding tasks transfer automatically to surviving nodes in the cluster so are still **available**
- If the failed node was the primary node then one of the remaining nodes is elected as the new primary node to continue managing HA
- Nodes can be added to a cluster to manage peak load, then later subtracted, without needing to redeploy Mule applications
- A cluster can be tuned to be either more performant or more reliable
 - The cluster can be configured with a quorum count so the cluster is not available until a minimum number of nodes are active
 - The cluster can be configured to decide how many nodes replicate the data, so data can survive if multiple nodes go offline

How customer-hosted Mule runtimes (nodes) join a cluster



- Nodes can discover and join a cluster using multicast or unicast

	Unicast	Multicast
Characteristics	<ul style="list-style-type: none">• Cluster uses IP address for identifying server	<ul style="list-style-type: none">• Cluster group servers automatically detect each other
Pros	<ul style="list-style-type: none">• No special network configuration other than IP of server	<ul style="list-style-type: none">• Nodes dynamically join the cluster when the node is started
Cons	<ul style="list-style-type: none">• IP of at least one other node must be known and configured in each node's cluster configuration	<ul style="list-style-type: none">• Only permitted in network where multicast is allowed

How shared memory is used in a cluster



- A cluster is implemented using Hazelcast to create a distributed **shared memory** data grid
 - Data is automatically replicated and available between the cluster's nodes
 - This allows data to survive if a node crashes or otherwise leaves the cluster
- Components that use a cluster's shared memory include
 - VM Queues
 - Object Stores
- Most connectors are not cluster-aware
 - But all connectors that use an Object Store are implicitly cluster-aware
 - Examples include: Cache scope, Idempotent Message validator, and the Round Robin router

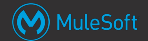
Features of clustering and load balancing



- Both clustering and load balancing in active-active scenarios have pros and cons

	Clustering	Load Balancing
Pros	<ul style="list-style-type: none">• Shared, distributed memory• Ideal for HA scenarios• Built-in load balancing for VM queues• Built into Mule	<ul style="list-style-type: none">• Easy to set up• No performance overhead due to latency or data replication• Configurable load balancing algorithms (round-robin, IP sticky, load-based, etc.)
Cons	<ul style="list-style-type: none">• Performance overhead due to latency and data replication• Not supported by CloudHub• Requires 3rd-party product to achieve HTTP load balancing	<ul style="list-style-type: none">• Requires third-party product• No data synchronization• Manage idempotency programmatically

Identify cluster aware connectors and design consideration for HA



- **Socket based**
 - Receives incoming traffic
 - Traffic must be distributed
 - Outbound socket based connectors don't need special consideration
 - Example: HTTP
- **Resource based**
 - Cluster automatically manages access to resource so only one clustered instance can access resource at a time
 - Outbound (writing) resource based connector generates unique resources
 - Examples: File, FTP
 - Distributed locking is not supported while writing

Comparing storage behavior for one or more customer-hosted Mule runtimes



	Non-persistent Object Store	Persistent Object Store
Standalone Mule runtime	<ul style="list-style-type: none">● in-memory store in the local Mule runtime● Data does not survive server restart	<ul style="list-style-type: none">● File-based store● Data survives server restart
Customer-hosted server group	<ul style="list-style-type: none">● Isolated in-memory store per Mule runtime● Data does not survive server restart	<ul style="list-style-type: none">● Isolated file-based store per Mule runtime● Data survives server restart
Cluster of customer-hosted Mule runtimes	<ul style="list-style-type: none">● in-memory store in the local Mule runtime● Data does not survive server restart	<ul style="list-style-type: none">● Hazelcast data-grid● The cluster itself can be configured to be only in-memory, or to use file-based storage

Identify cluster aware connectors and design consideration for HA

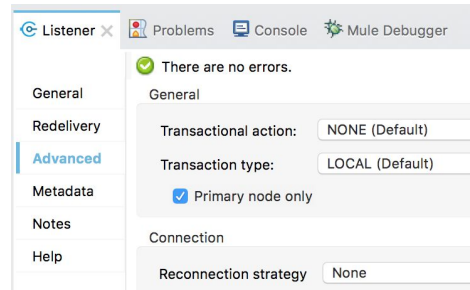


- Listener based

- Traffic is distributed automatically
- Must decide if the listener should only fire on the primary node or on all nodes
- Examples: VM, JMS

- Schedulers


- The scheduler only fires on the primary node



Load balancing for HTTP/S connector




- For customer-hosted Mule runtimes, HTTP requests need to be load balanced through a 3rd-party product
 - Traffic must be distributed
 - Load balancers is required (Nginx, Apache web server)
 - Outbound socket based connector don't need special consideration
- For CloudHub workers, HTTP requests are automatically load balanced through shared or dedicated CloudHub load balancers



Clustering for VM Connector

- Messages published to a VM queue in a cluster are automatically load balanced to receiving flows
 - No additional servers or infrastructure are required
 - Every node in the cluster can execute flows of deployed Mule apps
 - The cluster manager automatically determines what node to use based on load
 - Not a deterministic round-robin algorithm



All contents © Mule

19

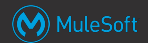
Understanding Anypoint MQ

- **Anypoint MQ** is a multi-tenant, cloud messaging service that enables customers to perform advanced asynchronous messaging scenarios between their applications
 - Is fully integrated with Anypoint Platform
 - Offers role based access control
 - Client application management
 - Connectors
 - Optionally, can provide strict first in, first out (FIFO) processing to enable ordering of messages
 - REST API allows access by non-Mule applications
 - Displays usage statistics on the number of messages

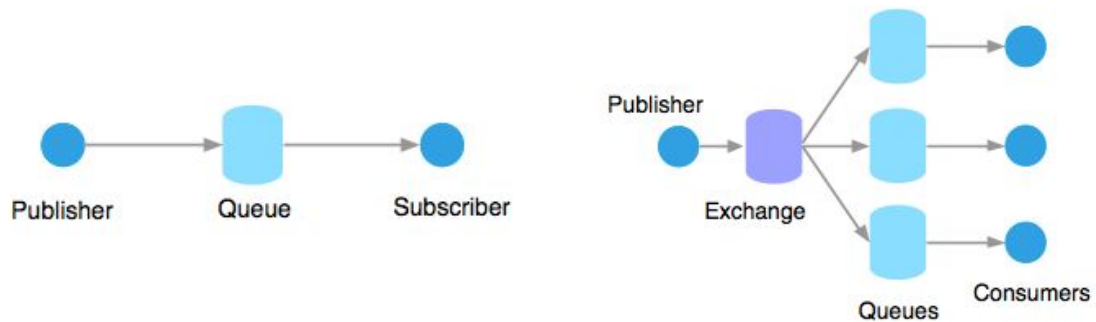
All contents © MuleSoft Inc.

20


Anypoint MQ message processing pattern



- An Anypoint MQ **queue** delivers a single message to a single consumer
- An Anypoint MQ **exchange** provides a way to broadcast a single message to many queues



How Anypoint MQ distributes messages



- In Anypoint MQ, reliability is provided via queues through a **lock and ack mode**
- A **consumer** retrieves messages from a **queue**
 - This locks the messages for a finite period of time, and makes the messages invisible to all other consumers
- If the message is not processed within the lock period, the lock times out, and consequently the message is made visible to other consumers to be processed
- This ensures that if there was a failure of some sort, such as the consumer node crashed, a message can be processed by another node

Messages in a VM queue are not load balanced for Mule apps deployed to a **standalone Mule runtime**



- Messages passed on a VM queue will NOT be automatically load balanced to receiving flows across Mule runtimes
 - Every standalone Mule runtime node will execute flow instances independently
 - The processing happens on a single node
 - No distributed processing

How VM connectors are load balanced in customer-hosted Mule runtimes



- When the Mule runtimes are combined into a cluster, persistent queues are backed by the memory grid
- Then multiple consumers share the VM queue, and messages are automatically load balanced between consumers
- The load balancing is as fast as possible, to whichever consumer is first (not round-robin)

- In **CloudHub**, each persistent VM queue is listened on by every **CloudHub worker**
 - But each message is read and processed **at least once** by only one CloudHub worker, and **duplicate processing is possible**
 - If the CloudHub worker fails, the message can be read by another worker to prevent loss of messages and this can lead to duplicate messaging
 - By default, every CloudHub worker's VM Listener receives different messages from the VM queue

- Unlike VM, JMS listener's **default behavior** is to **receive messages only in the *primary node***, no matter from what kind of destination messages are being consumed
- If consuming from a queue, the ***primary node*** configuration can be change to false to receive messages in all the nodes of the cluster
- Normal subscriptions where each subscriber will receive a copy of the published message, if consuming from topic with shared subscriptions mechanism (a mechanism for distributing messages to a set of subscribers to shared subscription topic), then you'll want to change the cluster configuration to consume messages only in the *primary node to false*

Load balancing for FTP, File connectors



- A cluster automatically manages access to a resource so only one clustered node can access the resource at a time
- An outbound resource based connector generates unique resources
- Distributed locking is not supported

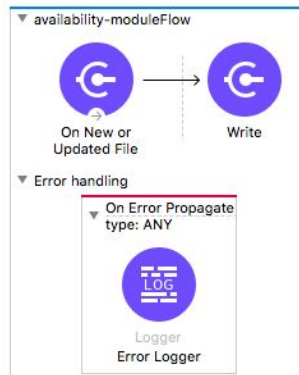
Failures of node in cluster / load balancer



- Failure of a node in a cluster or load balancer causes failures of HTTP requests processed by that node
 - However a typical enterprise level load balancer (such as AWS ELB) automatically recognizes failures of a node and stops sending subsequent requests to the failed node
- JMS, File, VM, FTP listeners, and Schedulers continue to perform processing on the elected primary node in cluster
 - The behavior is mostly consistent with the behavior in CloudHub
 - The primaryNodeOnly configuration is only used in a cluster, not by CloudHub
- Failures of a load balancer is a single point of failure
 - CloudHub load balancers are highly available and deployed in more than one availability zone

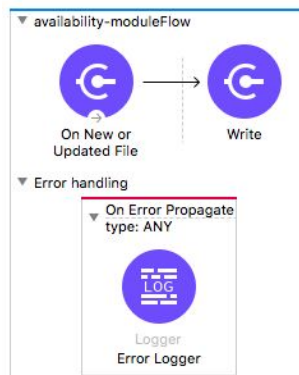
Exercise 13-1: Design an HA and reliable application for file transfer

- Design a file transfer use case that meets agreed HA goals, including minimum uptime and the minimum allowed number of Mule application failures



Exercise context

- The Mule application has to transfer a file from a local directory to an online FTP server
- The Mule application has to meet an HA availability SLA of 99.99% per year and must be highly reliable against Mule application failures

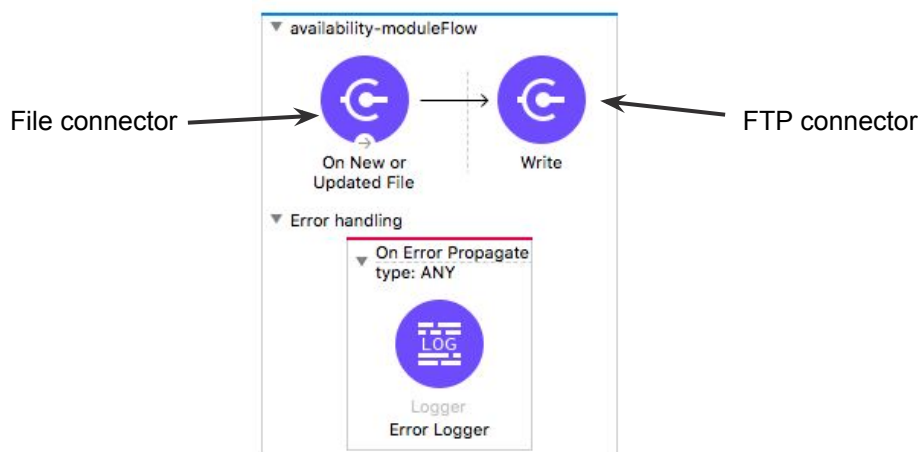


Exercise overview: Design an HA and reliable Mule application for file transfer

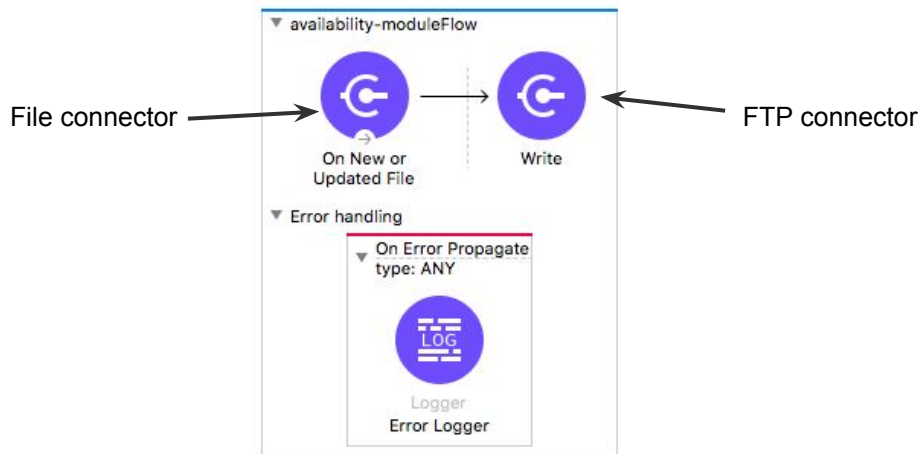
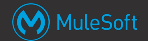


- Design flow/s for the Mule application
- Identify scaling options for the Mule application
- Identify the number of workers for the Mule application
- Decide Mule application persistence options to meet the Mule application's reliability goals
- Design the Mule application to meet reliability goals regarding Mule application and system crashes
- How would you design one flow to transfer a file to the FTP server
 - This is less reliable, but easier and more direct to implement

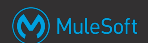
Exercise step: Design the file transfer flow



Exercise steps: What happens if the FTP write fails and the file is already deleted?



Exercise step: Answer these questions



- What failures points can reduce reliability of the system?
- If an error happens after reading the file, but before the file is completely processed, how can you assure the file is not deleted?
- What should be done if file reading fails a few times?
- What should be done if the FTP write process fails a few times?
- Can persistence help to enhance reliability?
- What are the options to maintain persistence?
- When should the file be deleted from the inbound file server?
- What horizontal or vertical scaling options can help?

Exercise solution options: Design an HA and reliable Mule application for file transfer

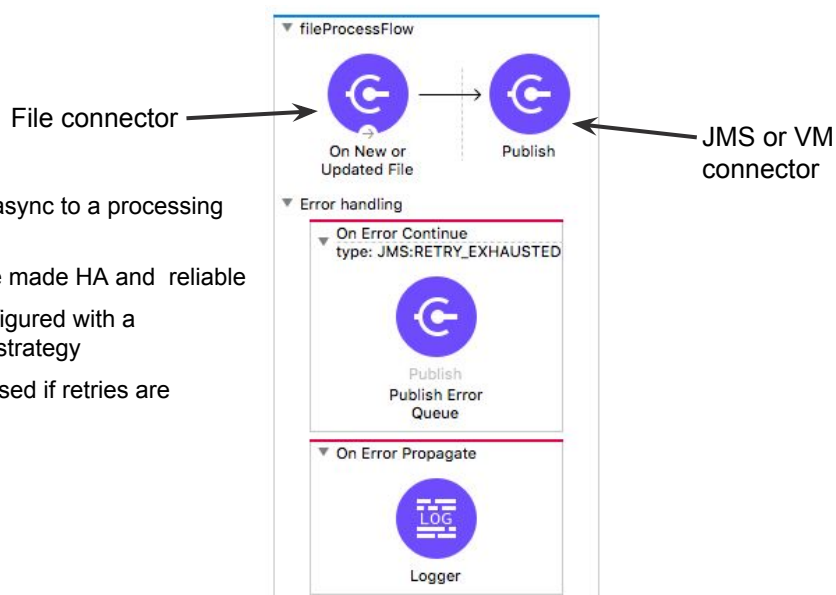


- Vertically scaling
 - A single Mule runtime does not make sense to use one node for HA (99.99%)
 - A better option is to horizontally scale the Mule application across multiple Mule runtimes
- Horizontal scaling options
 - Deploy to multiple Mule runtimes
 - Multiple nodes would access the same file
 - VM transport will not load balance between nodes
 - Clustering
 - Cluster manages resource access so only one node can access file
 - VM transport is selected for application persistence to enhance reliability in cluster

Exercise solution: How to orchestrate file transfer




- Send the file async to a processing queue
- Queue can be made HA and reliable
- Queue is configured with a reconnection strategy
- Error queue used if retries are exhausted



Exercise solution: How to orchestrate ftp process

JMS or VM
connector



FTP connector

- Receive files from the processing queue
- Queue can be made HA and reliable
- Queue is configured with a reconnection strategy
- For JMS, the Listener can also configure a redelivery policy
- Error queue used if JMS redeliveries are exhausted

- Write the processed file using an FTP connector
- Set a reconnection strategy
- If the Mule application is deployed to multiple Mule runtimes (or CloudHub workers), then several files can be processed concurrently

Summary



- HA goals must be clearly defined for applications
 - Performance and HA are often opposing goals, which involve trade-offs
- HA is achieved by scaling the Mule application with clustering and load balancing
- Load balancing and clustering for HTTP/S requests require an external product
- Automatic load balancing in a cluster only works for flows that start with a VM connector